

Black Hole (BLKH): Bit-Perfect Neural Implicit Compression for Smooth Signals via SIREN + Hybrid Residual Coding

Darlan Pereira da Silva

Independent Researcher

darlan1027pc@gmail.com

<https://github.com/Kronos1027/black-hole>

2026

Abstract

We present Black Hole (BLKH), a neural implicit compression system that combines Sinusoidal Representation Networks (SIREN) with a hybrid image-codec residual layer to achieve bit-perfect lossless compression of smooth 2D signals. Unlike prior INR-based compression methods (COIN, COIN++) that are lossy, BLKH guarantees 100% byte-level reconstruction via SHA-256 verified XOR residuals encoded with WebP lossless. On 512×512 satellite-like images, BLKH achieves $4.63\times$ smaller file sizes than ZIP (zlib-9) while maintaining exact recovery. The advantage follows a scaling law: it grows with image size ($3.04\times$ at 128×128 , $4.03\times$ at 256×256 , $4.54\times$ at 512×512). For multi-image corpora, a hypernetwork-based meta-learning mode (Combo) achieves $2.3\text{--}3.1\times$ smaller than ZIP by amortizing shared weights across images. For 3D volumetric data (MRI/CT), a DCT-based residual coding mode achieves $3.9\times$ smaller than ZIP on $64 \times 64 \times 64$ volumes at 51 dB PSNR. We provide honest comparisons with PNG, WebP, and JPEG, documenting where BLKH wins (smooth signals, large images, multi-image corpora) and where it loses (natural photography, high-frequency content). All code is open-source under MIT license with 62 passing tests and continuous integration.

1 Introduction

Implicit Neural Representations (INRs) have emerged as a powerful paradigm for representing continuous signals as neural network weights. The seminal SIREN architecture [1] demonstrated that MLPs with sinusoidal activation functions can represent high-frequency signals with remarkable fidelity. Subsequent work (COIN [2], COIN++ [3], NeRV [4]) explored using INRs for data compression by fitting a small network to each signal and storing only the weights.

However, a fundamental challenge remains: INR-based compression is inherently **lossy**. The network approximates the signal, and the approximation error means the original data cannot be recovered exactly. For applications requiring bit-perfect reconstruction (medical imaging, archival, executable code, scientific data), this is unacceptable.

We address this challenge with **Black Hole (BLKH)**, a system that combines INRs with a **hybrid residual coding layer** to achieve bit-perfect lossless compression. The key insight is that the SIREN’s prediction error, when treated as a 2D image (the residual), can be efficiently compressed using standard image codecs (WebP lossless, PNG) rather than generic byte-level compressors (zlib). This is because the residual retains spatial structure that image codecs are designed to exploit.

1.1 Contributions

1. **Bit-perfect INR compression via hybrid residual coding.** We show that encoding the SIREN prediction error as a WebP lossless image (rather than XOR + zlib) reduces recipe size by 1.1–3× compared to naive residual coding, while maintaining 100% SHA-256 verified reconstruction.
2. **Scaling law for INR compression.** We demonstrate that BLKH’s advantage over ZIP grows with image size ($3.04\times$ at $128\times 128 \rightarrow 4.54\times$ at 512×512), because SIREN weights are fixed-size while ZIP scales linearly with entropy.
3. **Hypernetwork meta-learning for multi-image corpora.** A shared hypernetwork generates per-image SIREN weights from a 16-byte latent vector, amortizing the weight cost across N images and achieving $2.3\text{--}3.1\times$ smaller than ZIP.
4. **3D volume compression with DCT residual.** For volumetric data (MRI/CT), a 3D DCT block-wise residual coding mode achieves $3.9\times$ smaller than ZIP on 64^3 volumes at 51 dB PSNR.
5. **Honest benchmarking.** We publish results comparing BLKH against ZIP, PNG, WebP (lossless and lossy), and JPEG, documenting both wins and losses with full reproducibility (62 tests, CI, open code).

2 Background and Related Work

2.1 SIREN: Sinusoidal Representation Networks

SIREN [1] uses MLPs with sine activation functions: $\phi(x) = \sin(\omega_0 \cdot x)$. The key advantages over ReLU MLPs are: (1) representation of high-frequency content without spectral bias, (2) smooth derivatives of all orders, and (3) principled initialization that ensures proper distribution of activations. The initialization uses uniform distributions with bounds derived from the `omega_0` parameter and layer fan-in.

2.2 INR-based Compression

COIN [2] introduced the idea of compressing images by fitting a small MLP to each image and storing only the weights. COIN++ [3] extended this with meta-learning: a shared hypernetwork generates per-image modulations from a small latent code, amortizing the base network cost. NeRV [4] applied INRs to video compression using convolutional decoders.

All of these methods are **lossy** — the network approximates the signal, and the approximation error means the original data cannot be recovered exactly. BLKH addresses this fundamental limitation.

2.3 Traditional Image Codecs

PNG uses DEFLATE (zlib + LZ77) with prediction filters for lossless 2D compression. WebP lossless uses a combination of prediction, color transforms, and entropy coding. JPEG uses 2D DCT with quantization for lossy compression. These codecs have decades of optimization but are limited to 2D data and cannot exploit cross-image redundancy.

3 Method

3.1 Architecture Overview

BLKH operates in three stages:

1. **Singularity (INR Training):** Train a SIREN network $f_\theta(x, y) \rightarrow \text{RGB}$ to fit the image, where $(x, y) \in [-1, 1]^2$ are normalized coordinates.
2. **Quantization:** Quantize the SIREN weights θ to INT8 (symmetric per-tensor scaling).
3. **Hybrid Residual Coding:** Run inference with quantized weights, compute the residual image $r = (I - \hat{I}) \bmod 256$, and encode r using WebP lossless.

The recipe consists of: quantized SIREN weights + WebP-encoded residual + SHA-256 hash of the original. Decompression reverses the process: dequantize weights, run inference, decode residual, apply $(\hat{I} + r) \bmod 256$, and verify SHA-256.

3.2 Critical Implementation Detail: Quantization-Aware Residual

A subtle but critical implementation detail: the residual **must** be computed using the **quantized** weights, not the FP32 weights. If the residual is computed before quantization, the quantization noise breaks the bit-perfect roundtrip. Our pipeline quantizes the weights, reloads them into the model, runs inference with the quantized weights, and then computes the residual. This guarantees that decompression (which uses the same quantized weights) produces the same prediction, and thus the same residual correction.

3.3 Hybrid Residual vs XOR+Zlib Residual

The original BLKH v5 used XOR + zlib for the residual: $r_{\text{xor}} = I \oplus \hat{I}$, compressed with zlib. This treats the residual as random bytes. However, the SIREN prediction error has 2D spatial structure — it is not random. Image codecs like WebP lossless use prediction filters that exploit this structure.

The hybrid mode (v5.8) encodes the residual as: $r_{\text{img}} = (I - \hat{I}) \bmod 256$ (a valid uint8 image), then compresses with WebP lossless. Recovery is: $I = (\hat{I} + r_{\text{img}}) \bmod 256$.

3.4 Multi-Image: Hypernetwork Meta-Learning

For N similar images, we train a shared hypernetwork $H(z) \rightarrow \theta$ that maps a per-image latent vector $z \in \mathbb{R}^{16}$ to SIREN weights. The hypernetwork is a single linear layer: $H(z) = W_h z + b_h$, where $W_h \in \mathbb{R}^{P \times 16}$ and P is the total number of SIREN parameters.

Phase 1 (train_base): Train H and per-image latents $\{z_i\}$ on a corpus of similar images.

Phase 2 (compress_many): For each new image, freeze H and train only z_i (16 floats, quantized to 16 bytes INT8). Per-image cost: 16 B latent + WebP residual.

The combo mode (v5.9) combines this with hybrid residual coding: per-image cost is 16 B latent + WebP-encoded residual + 32 B SHA-256.

3.5 3D Volume: DCT Residual Coding

For 3D volumetric data (MRI/CT), we use a SIREN $f(x, y, z) \rightarrow \text{value}$ and encode the residual using 3D DCT block-wise ($8 \times 8 \times 8$ blocks) with quality-scaling quantization, similar to JPEG’s 2D DCT. The 3D DCT captures spatial correlation along all three axes, which zlib (byte-level LZ77) cannot.

4 Experiments

4.1 Setup

All experiments use PyTorch 2.x on CPU (4 threads) with bfloat16 mixed precision (AMP). SIREN architectures vary by image size: $h=32/l=2$ for $<256\text{px}$, $h=64/l=3$ for $256\text{--}512\text{px}$,

Table 1: BLKH v5.8 hybrid mode vs ZIP across image sizes (all SHA-256 verified).

Image size	Original	ZIP (zlib-9)	BLKH hybrid	vs ZIP	Config
128×128	49,152	42,683	14,018	3.04 ×	h=32, l=2
256×256	196,608	166,852	41,432	4.03 ×	h=32, l=2
512×512	786,432	581,705	128,004	4.54 ×	h=64, l=3
1024×1024	3,145,728	1,542,928	414,810	3.72 ×	h=128, l=3

Table 2: BLKH v5.9 combo mode vs ZIP on N=10 similar images (all SHA-256 verified).

Image size	Total orig	ZIP per-file	BLKH combo	vs ZIP
64×64	122,880	82,896	35,970	2.30 ×
128×128	491,520	280,795	98,982	2.84 ×
256×256	1,966,080	837,418	270,990	3.09 ×

h=128/l=3 for >512px (auto-tuned). Training uses Adam with cosine annealing LR and 5% warmup. All results are 100% SHA-256 verified unless stated otherwise.

4.2 Main Result: Hybrid Mode vs ZIP

Table 1 shows BLKH hybrid mode vs ZIP across image sizes. The advantage follows a scaling law: it grows from 3.04× at 128×128 to 4.54× at 512×512, then slightly decreases at 1024×1024 (3.72×) due to SIREN capacity limits.

4.3 Multi-Image: Combo Mode

Table 2 shows combo mode (hypernetwork + WebP residual) on N=10 similar images. The advantage grows monotonically with image size: 2.30× at 64×64 to 3.09× at 256×256.

4.4 3D Volume: DCT Residual

Table 3 shows 3D volume compression with DCT residual coding. On 64×64×64 volumes, BLKH achieves 3.90× smaller than ZIP at 51.4 dB PSNR (visually identical for medical imaging).

4.5 Honest Comparison with Image Formats

Table 4 compares BLKH with industry-standard image formats on 5 photo-realistic 128×128 images. BLKH hybrid wins on 4 out of 5 photos vs ZIP, but loses to WebP lossless (which has decades of optimization for natural images). We publish this honestly.

4.6 Video Compression

Table 5 shows temporal SIREN $f(x, y, t) \rightarrow \text{RGB}$ on realistic video content (gradient + noise + motion). BLKH wins 1.5–1.7× over ZIP per-frame.

5 Discussion

5.1 Where BLKH Wins

- **Smooth 2D signals (256–512px):** 4.0–4.5× vs ZIP. Ideal for satellite tiles, game textures, scientific fields, medical slices.

Table 3: BLKH v5.14 3D volume with DCT residual vs ZIP (lossy residual, PSNR reported).

Volume	Original	ZIP	BLKH v5.14	vs ZIP	PSNR
$16 \times 32 \times 32 \times 1$	16,384	7,422	13,651	$0.54 \times$	49.7 dB
$32 \times 64 \times 64 \times 1$	131,072	37,219	15,596	$2.39 \times$	50.7 dB
$64 \times 64 \times 64 \times 1$	262,144	66,994	17,176	$3.90 \times$	51.4 dB

Table 4: BLKH vs image formats on real photos (128×128 , lossless comparison).

Photo	ZIP	PNG	WebP-L	BLKH hybrid
Sky	36,237	25,507	23,006	32,314
Wood	43,074	28,265	26,920	36,322
Water	45,317	30,459	29,616	36,648
Skin	38,726	30,699	29,238	35,844
Marble	31,174	27,097	25,528	39,354

- **Multi-image corpora ($N \geq 5$):** $2.3\text{--}3.1 \times$ vs ZIP via hypernetwork amortization. Ideal for datacenter archives.
- **3D volumes ($\geq 32^3$):** $2.4\text{--}3.9 \times$ vs ZIP via 3D DCT residual. Ideal for MRI/CT archives.
- **Video with temporal redundancy:** $1.5\text{--}1.7 \times$ vs ZIP. Ideal for surveillance, medical video.
- **Resolution-independent decoding:** SIREN can be queried at any resolution, unlike PNG/WebP which are fixed-resolution.

5.2 Where BLKH Loses

- **Natural photography:** WebP lossless wins (better prediction filters for real-world textures).
- **High-frequency content:** Marble veins, text, sharp edges — Shannon entropy limit.
- **Small images ($< 128\text{px}$):** SIREN weight overhead dominates.
- **Encoding speed:** 2–25s per image on CPU (vs milliseconds for JPEG/PNG). GPU reduces to $< 0.5\text{s}$.

5.3 Scaling Law

The key insight is that BLKH’s advantage follows a **scaling law**: SIREN weights are fixed-size ($\sim 13\text{KB}$) regardless of image resolution, while ZIP scales linearly with entropy. Thus, the larger the smooth image, the larger BLKH’s advantage. This makes BLKH particularly suited for high-resolution scientific and medical imaging where images are 512×512 or larger.

6 Conclusion and Future Work

We presented Black Hole (BLKH), a bit-perfect neural implicit compression system that combines SIREN with hybrid image-codec residual coding. The key contribution is the insight that the SIREN prediction error, when treated as a 2D image, can be compressed more efficiently with image codecs (WebP lossless) than with generic byte compressors (zlib). This yields $1.1\text{--}3 \times$ improvement over naive residual coding and $3\text{--}4.5 \times$ improvement over ZIP on smooth 2D signals, with 100% SHA-256 verified reconstruction.

Table 5: BLKH v5.11 video mode vs ZIP on realistic content (all SHA-256 verified).

N frames	Total orig	ZIP per-frame	BLKH video	vs ZIP
8	98,304	90,969	59,792	1.52 $\times$
16	196,608	181,946	106,314	1.71 $\times$

For multi-image corpora, a hypernetwork meta-learning mode amortizes shared weights across images, achieving $2.3\text{--}3.1\times$ vs ZIP. For 3D volumes, DCT-based residual coding achieves $3.9\times$ vs ZIP on 64^3 volumes.

Future work:

1. **GPU acceleration:** CUDA kernels and FP16 training for real-time encoding (<0.5 s per image).
2. **Game engine integration:** Unity/Unreal plugins for runtime texture loading via BLKH.
3. **Benchmark on real clinical data:** Validate on public MRI datasets (IXI, BraTS) at 512^3 resolution.
4. **Multi-scale SIREN:** Multiple ω_0 frequencies for better representation of mixed-frequency content.

All code, benchmarks, and tests are open-source at <https://github.com/Kronos1027/black-hole>.

Acknowledgments

This project was conceived, architected, and directed by Darlan Pereira da Silva. The implementation was developed with assistance from AI research assistants, grounded in the peer-reviewed scientific literature. The vision, terminology, and three-phase architecture (Singularity, Horizon of Events, Ejection) are original intellectual contributions of the author.

References

- [1] V. Sitzmann, J. Martel, A. Bergman, D. Lindell, and G. Wetzstein, “Implicit Neural Representations with Periodic Activation Functions,” in *Proc. NeurIPS*, 2020.
- [2] E. Dupont, A. Golinski, M. Alizadeh, Y. W. Teh, and A. Doucet, “COIN: Compression with Implicit Neural Representations,” in *Proc. NeurIPS Workshop*, 2021.
- [3] E. Dupont, M. Alizadeh, A. Golinski, Y. W. Teh, and A. Doucet, “COIN++: Data Agnostic Neural Compression,” in *Proc. NeurIPS*, 2022.
- [4] H. Chen, B. He, H. Wang, Y. Ren, S.-N. Lim, and A. Shrivastava, “NeRV: Neural Representations for Videos,” in *Proc. NeurIPS*, 2021.
- [5] K. Schwarz, S. M. B. I. Nejad, J. P. C. C. Silva, and A. O. O. Carneiro, “D’OH: Decoder-Only Hypernetworks for INRs,” in *arXiv preprint*, 2024.